

# VibeWalk: AI-Powered Safety Navigation System

---

## System Design & Technical Report (Convolve 4.0 Submission)

**Project Name:** VibeWalk

**Submission Category:** Semantic Search & Vector Memory Systems

**Primary Engine:** Qdrant Vector Database

---

### 1. Problem Statement

#### 1.1 Societal Issue: The Safety Gap in Navigation

Conventional navigation apps (Google Maps, Waze) optimize primarily for **speed** and **efficiency**. They calculate the shortest path from A to B, often routing pedestrians through unlit alleyways, high-crime zones, or construction sites without warning.

For vulnerable groups **women, children, and the elderly** the "fastest" route is not always the best route. "Safety" is subjective and hyper local; a street that feels safe at 2 PM might be terrifying at 10 PM.

#### 1.2 Impact

- **Fear of Movement:** 45% of women report avoiding walking alone at night due to safety concerns.
- **Information Asymmetry:** Crime data exists (e.g., NYC Open Data), but it is buried in CSV files, not integrated into daily navigation tools.
- **Static Data Failure:** Official crime statistics are lagging indicators. They don't capture real-time temporary hazards like a broken streetlight, a gathered crowd, or harassment incidents that go unreported to police.

#### 1.3 The Solution: VibeWalk

VibeWalk is a "**Safety-First**" **Semantic Navigation System**. Instead of optimizing for distance, it optimizes for "**Vibe**" (**Semantic Safety Score**). It uses **Qdrant** to store and retrieve dense vector representations of the city combining hard crime stats, positive place reviews, and real-time user reports to route users through the safest, most comfortable paths.





Figure 1: VibeWalk generating safe routes using semantic scoring.

## 2. System Design

### 2.1 Architecture Overview

The system follows a modern **RAG (Retrieval-Augmented Generation)** pattern applied to Navigation.



### 2.2 Why Qdrant is Critical

Qdrant is not just a database for VibeWalk; it is the **Decision Engine**.




Figure 2: Qdrant retrieving "Semantic Intel" (Reviews) near a user.

1. **Semantic Flexibility:** We don't just filter by "crime\_count > 5". We compare the *semantic meaning* of reports. A report of "loud music" implies a different safety profile than "armed robbery". Vector search understands this nuance.
2. **Geo-Spatial Vector Search:** Qdrant's unique ability to combine **Dense Vector Search** (semantic similarity) with **Geo-Radius Filtering** in a single query is the backbone of our scoring algorithm.
3. **Real-Time "Evolving Memory":** User reports are immediately vectorized and upserted. Qdrant's speed allows these new memories to influence route scores instantly (within milliseconds).

### 2.3 Database Structure (Payload Schema)

Each "Vibe Node" in Qdrant represents a semantic data point in the city.

Field	Type	Description	Example
<code>id</code>	UUID	Unique Identifier	550e8400-e29b...
<code>vector</code>	<code>float[384]</code>	BGE-Small-EN Embedding	[0.02, -0.15, ...]
<code>payload.text</code>	String	Description of event/place	"Dark alley, saw a fight."
<code>payload.type</code>	String	Source Category	crime, review, user_report
<code>payload.location</code>	GeoPoint	Lat/Lng	{ lat: 40.75, lon: -73.99 }
<code>payload.timestamp</code>	ISO8601	Time relevance	2023-10-27T10:00:00Z
<code>payload.severity</code>	String	Impact Level	high, medium, low

### 3. Multimodal Strategy

VibeWalk treats the city as a "Multimodal Text". We convert distinct signals into a unified vector space.

### 3.1 Data Types & Sources

1. **Structured Government Data (Textual)**: NYC Socrata Crime API.
  - *Raw*: "PD Code 101, Assault 3"
  - *Vectorized*: "Violent assault crime report"
2. **Unstructured Social Proof (Textual)**: Place Reviews.
  - *Raw*: "Lovely park with security guards visible."
  - *Vectorized*: "Safe, monitored, family friendly location."
3. **Real-Time Crowd Intelligence (User inputs)**:
  - *Raw*: "Street light broken here, very dark."
  - *Vectorized*: "Low visibility, potential hazard, fear."

### 3.2 Embedding Strategy

We use **FastEmbed** ([BAAI/bge-small-en-v1.5](#)) running locally on the backend. This model is optimized for semantic retrieval.

- **Concept Anchors**: We pre-compute "Concept Vectors" for *Safety* and *Danger*.
    - **DANGER\_CONCEPT** = Embedding("Crime, assault, robbery, darkness, fear")
    - **SAFE\_CONCEPT** = Embedding("Happy, families, well-lit, security, police")
- 

## 4. Search, Memory, & Recommendation Logic

### 4.1 Search: The "Route Scoring" Algorithm

How do we turn vectors into a "7/10 Safety Score"?

1. **Path Generation**: Fetch 3 diverse physical paths from OSRM (Direct, North-Alt, South-Alt).
2. **Sampling**: We sample equidistant points (e.g., every 50 meters) along each path.
3. **Semantic Probe**: For each point, we query Qdrant:

```
# Pseudocode
hits = qdrant.search(
    collection="nyc_vibes",
    query_vector=DANGER_CONCEPT,
    filter=GeoRadius(center=point, radius=100m)
)
```

#### 4. Scoring:

- If **similarity\_score > 0.60** (High Relevance Match):
- Apply **Penalty**: **Score -= Hit.Score \* 4.0**
- The penalty is *semantic*: A match for "murder" (Similarity 0.85) penalizes more than "littering" (Similarity 0.40).

### 4.2 Memory: Evolving Safety Map

When a user reports a hazard, the system "remembers":

1. **Ingest:** POST `/report` receives text ("Creepy van parked here").
2. **Vectorize:** Convert text to vector.
3. **Upsert:** Store in Qdrant with `source="user_report"` and `timestamp=NOW`.




Figure 3: User reporting a hazard, which becomes a permanent vector memory. 4. **Recall:** The very next route search query will include this new vector in its search radius.




Figure 4: The system recalling a user report as part of its nearby intelligence.

If semantically dangerous, it immediately lowers the safety score of that road segment for *all other users*.

#### 4.3 Recommendation Logic

We don't just warn; we guide.

- **Safe Havens:** We query for `SAFE_CONCEPT` (Reviews, Open Businesses).
- **Contextual Display:** If a route is dangerous, the frontend suggests "Safe Havens" (e.g., "Starbucks (Open)", "Police Station") along the path as waypoints.

## 5. Limitations & Ethics

### 5.1 Algorithmic Bias

- **Risk:** If "crime" data is historically biased against certain neighborhoods, the AI might flag minority neighborhoods as "unsafe" unfairly.
- **Mitigation:** We mix in **Positive Reviews** (Social Proof) to counterbalance raw crime stats. A busy street in a "rough" neighborhood might be safer than a deserted street in a "rich" one due to "eyes on the street" (Jacobs' efficiency).

### 5.2 Failure Modes

- **Sparse Data:** In areas with no reports or reviews, the system defaults to "Safe" (10/10), which is a false negative risk.
- **Adversarial Attacks:** Bad actors could spam false "danger" reports to reroute traffic.
  - *Future Fix:* Reputation systems and outlier detection using clustering mechanisms in Qdrant.

### 5.3 Privacy

- User location data is transient (sent for routing) and not stored.
  - User reports are anonymized before vectorization.
- 

## 6. Future Roadmap

### 6.1 Visual Evidence & Photo Verification

Currently, VibeWalk relies on text. Future versions will allow users to **attach photos** to their reports (e.g., specific images of broken infrastructure or dangerous conditions).

- **Multimodal Learning:** We will integrate **CLIP-based models** to "read" these images. If a user uploads a photo of a dark alley, the system will automatically extract vector concepts like **darkness**, **isolation**, and **decay** without the user needing to type a description.
- **Trust & Verification:** Photos provide a "proof of work" to combat fake reviews. A report of "graffiti" accompanied by a verified photo is weighted higher than text alone.

### 6.2 Advanced Safety Features

- **IoT Integration:** Connecting with smart city APIs to ingest real-time data from connected streetlights and noise sensors.
  - **SOS & Emergency Mode:** A "Panic Mode" that instantly shares the user's location with trusted contacts and routes them to the nearest "High Safety Score" safe haven (e.g., Police Station, 24/7 Store).
  - **User Reputation System:** A gamified trust score for reporters. Users with a history of verified, helpful reports will have their inputs weighted more heavily in the safety algorithm.
- 

## 7. Conclusion

VibeWalk demonstrates that **Vector Search** is more than just a "Create-Read-Update-Delete" (CRUD) tool. It is a logical reasoning engine. By encoding the *qualitative* feeling of safety into high-dimensional space, we

provide a tool that protects the most vulnerable members of society, translating "Vibes" into actionable, life saving navigation decisions.